


Hough Transform



Outline

- Hough transform
- Homography

Voting schemes

- Let each feature vote for all the models that are compatible with it
- Hopefully the noise features will not vote consistently for any single model
- Missing data doesn't matter as long as there are enough features remaining to agree on a good model

The Hough Transform

- A powerful method for detecting curves from boundary information.
- Exploits the duality between points on a curve and parameters of the curve.
- Can detect analytic as well as non-analytic curves

Hough transform

- An early type of voting scheme
- General outline:
 - Discretize parameter space **into bins**
 - For each **feature point in the image**, put a **vote** in every **bin in the parameter space** that could have generated this point
 - Find bins that have the **most votes**







Image space



Hough parameter space

Analytic representation of a line


- ▶ In the analytic representation of a line $y = mx + c$, every choice of parameters (m, c) represents a different line.
- ▶ This is known as the *slope-intercept* parameter space.
- ▶ Weakness: vertical lines have $m = \infty$.




Parameter space representation

- A line in the image corresponds to a point in Hough space

Image space




Hough parameter space




Parameter space representation

- What does a point (x_0, y_0) in the image space map to in the Hough space?




Parameter space representation

- What does a point (x_0, y_0) in the image space map to in the Hough space?
 - Answer: the solutions of $c = -x_0m + y_0$
 - This is a line in Hough space



Parameter space representation


- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?




Parameter space representation

- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?
 - It is the intersection of the lines $c = -x_0m + y_0$ and $c = -x_1m + y_1$

Image space




Hough parameter space



Parameter space representation

- Problems with the (m,c) space:
 - Unbounded parameter domain
 - Vertical lines require infinite m
- Alternative: polar representation



$$r = x \cos(\theta) + y \sin(\theta)$$

Each point will add a sinusoid in the (θ, ρ) parameter space

Polar representation of a line

- Solution: Polar representation (r, θ) where
 - r = perpendicular distance of line from origin
 - θ = angle of vector orthogonal to the line
- Every (r, θ) pair represents a 2D line.

Additional formulae

$$\sin(A + B) = \sin A \cos B + \sin B \cos A$$


$$\sin(A - B) = \sin A \cos B - \sin B \cos A$$

$$\cos(A + B) = \cos A \cos B - \sin A \sin B$$

$$\cos(A - B) = \cos A \cos B + \sin A \sin B$$

$$\tan(A + B) = \frac{\tan A + \tan B}{1 - \tan A \tan B}$$

$$\tan(A - B) = \frac{\tan A - \tan B}{1 + \tan A \tan B}$$



$$y = mx + c$$

$$m = \tan(\alpha) = \tan\left(\theta + \frac{\pi}{2}\right)$$

$$= \frac{\sin\left(\theta + \frac{\pi}{2}\right)}{\cos\left(\theta + \frac{\pi}{2}\right)} = \frac{\cos(\theta)}{-\sin(\theta)}$$

$$c = \frac{r}{\sin(\theta)}$$

$$y = -\frac{\cos(\theta)}{\sin(\theta)}x + \frac{r}{\sin(\theta)}$$

$$r = x \cos(\theta) + y \sin(\theta)$$

Hough Transform for Line Detection

- ▶ An algorithm for finding lines given some edge points.
- ▶ Given point (x, y) , line passing through it with angle θ must have perpendicular $r = x \cos(\theta) + y \sin(\theta)$.
- ▶ Given any edge pixel (x, y) , potentially 180 lines could pass through it assuming angular resolution of 1° .
- ▶ Looping through the angles gives (r, θ) pairs for all lines through (x, y) .
- ▶ So pixel (x, y) should *vote for* all those lines.




Figure: Lines passing through a point. **Left:** Angular resolution of 30° . **Right:** Angular resolution of 10° . Author: N. Khan (2021)

Hough Transform for Line Detection




Figure: The accumulator array used to gather votes for each line. Each (r, θ) pair needs to be quantized into bin-indices before casting a vote. Author: N. Khan (2019)

Hough Transform for Line Detection

- By repeating this process for all edge pixels, actual lines will get a high number of votes.




Figure: Each point votes for every line that passes through it. Genuine lines will get more votes. Author: N. Khan (2021)

Hough Transform for Line Detection

- Pseudocode

initialize 2D (vote) accumulator array A to all zeros.

for every edge point (x, y)

for $\theta = 0$ to π

compute $r = x \cos(\theta) + y \sin(\theta)$

compute indices $(r_{\text{ind}}, \theta_{\text{ind}})$ corresponding to (r, θ)

increment $A(r_{\text{ind}}, \theta_{\text{ind}})$ by 1 \leftarrow vote of point (x, y) for line (r, θ)

valid lines are where $A > \text{threshold}$

Hough Transform for Line Detection (detailed Pseudocode)

1. $\theta_{\text{range}} = 180^\circ$
2. $\theta_{\text{binsize}} = 1^\circ$ (for example)
3. $\theta_{\text{size}} = \left\lceil \frac{\theta_{\text{range}}}{\theta_{\text{binsize}}} \right\rceil$
4. $r_{\text{max}} = \text{length of image diagonal}$
5. $r_{\text{range}} = 2r_{\text{max}}$
6. $r_{\text{binsize}} = 1$ pixel (for example)
7. $r_{\text{size}} = \left\lceil \frac{r_{\text{range}}}{r_{\text{binsize}}} \right\rceil$
8. initialize 2D (vote) accumulator array A of size $(r_{\text{size}}, \theta_{\text{size}})$ to all zeros.
9. for every edge point (x, y)
10. for $\theta = 0$ to θ_{range}
11. compute $r = x \cos(\theta) + y \sin(\theta)$
12. $r_{\text{ind}} = \text{round} \left(\frac{r + r_{\text{max}}}{r_{\text{binsize}}} \right)$
13. $\theta_{\text{ind}} = \text{round} \left(\frac{\theta \bmod 180}{\theta_{\text{binsize}}} \right)$
14. increment $A(r_{\text{ind}}, \theta_{\text{ind}})$ by 1 \leftarrow vote of point (x, y) for line (r, θ)

Hough Transform for Line Detection (detailed Pseudocode)

15. smooth votes via Gaussian convolution with kernel G_{σ_h} to account for uncertainties in the gradient direction
16. perform non-maxima suppression in $k \times k$ neighborhoods to remove fake lines around real ones
17. valid lines¹ are where $A > \tau$ which can be computed as a percentile

Improvement


After edge detection, we already know the gradient direction at (x, y) .

- ▶ So there is no need to iterate over all possible θ .
- ▶ Use the correct θ from the gradient direction.
- ▶ This removes the loop at step 10.
- ▶ Pixel (x, y) only votes for the line that was actually passing through it.

This speeds-up the algorithm.


Extension: Incorporating image gradients

- Recall: when we detect an edge point, we also know its gradient direction
- But this means that the line is uniquely determined!
- Modified Hough transform:
 - For each edge point (x,y)
 - θ = gradient orientation at (x,y)
 - $r = x \cos \theta + y \sin \theta$
 - $H(\theta, r) = H(\theta, r) + 1$
- end




$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

Basic illustration




features




votes

Other shapes



Square




Circle




Several lines




Effect of noise



Effect of noise



features




votes



- Peak gets fuzzy and hard to locate

Effect of noise

- Number of votes for a line of 20 points with increasing noise:



Random points




features

votes

- Uniform noise can lead to spurious peaks in the array

Random points

- As the level of uniform noise increases, the maximum number of votes increases too:




Hough transform: Pros

- Can deal with non-locality and occlusion
- Can detect multiple instances of a model in a single pass
- Some robustness to noise: noise points unlikely to contribute consistently to any single bin

Hough transform: Cons

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- It's hard to pick a good grid size

Results




Original



Using edge pixels only

$\tau = 95$ -th percentile



Using edge pixels and
gradient orientations


$\tau = 70$ -th percentile

Figure: Line detection via Hough transform. Canny parameters: $\sigma_e = 1$, $t_h = 80$ -th percentile, $t_l = 40$ -th percentile. Hough parameters: $\sigma_h = \frac{\sigma_e}{5}$, $k = 3$. Author: N. Khan (2021)

Results




Original



Using edge pixels only

$\tau = 95$ -th percentile



Using edge pixels and
gradient orientations

$\tau = 70$ -th percentile

Figure: Line detection via Hough transform. Canny parameters: $\sigma_e = 1$, $t_h = 80$ -th percentile, $t_l = 40$ -th percentile. Hough parameters: $\sigma_h = \frac{\sigma_e}{5}$, $k = 3$. Author: N. Khan (2021)

Results




Original



Using edge pixels only

$\tau = 95$ -th percentile



Using edge pixels and
gradient orientations

$\tau = 90$ -th percentile

Figure: Line detection via Hough transform. Canny parameters: $\sigma_e = 1$, $t_h = 80$ -th percentile, $t_l = 40$ -th percentile. Hough parameters: $\sigma_h = \frac{\sigma_e}{5}$, $k = 3$. Author: N. Khan (2021)

Hough transform for circles

- How many dimensions will the parameter space have?
- Given an oriented edge point, what are all possible bins that it can vote for?

Hough Transform for Circle Detection

- ▶ Analytic representation of circle of radius r centered at (a, b) is $(x - a)^2 + (y - b)^2 - r^2 = 0$.
- ▶ Hough space has 3 parameters (a, b, r) .

Pseudocode

For every boundary point (x, y)

 For every (a, b) in image plane

 Compute $r(a, b) = \sqrt{(x - a)^2 + (y - b)^2}$

 Compute $a_{\text{ind}}, b_{\text{ind}}$ and r_{ind}

 Increment $A(a_{\text{ind}}, b_{\text{ind}}, r_{\text{ind}})$ by 1

NMS($A * G_{\sigma_h}$) $> \tau$ represents valid circles.

- ▶ If we know the gradient vector $\nabla I(x, y)$ at point (x, y) , then we also know that the center (a, b) can only lie along this line.
- ▶ Hough space still has 3 parameters (a, b, r) but we search for r over a 1D space instead of a 2D plane.

Pseudocode

For every boundary point (x, y)

For every (a, b) **along gradient vector $\nabla I(x, y)$**

Compute $r(a, b) = \sqrt{(x - a)^2 + (y - b)^2}$

Compute $a_{\text{ind}}, b_{\text{ind}}$ and r_{ind}


Increment $A(a_{\text{ind}}, b_{\text{ind}}, r_{\text{ind}})$ by 1

$\text{NMS}(A * G_{\sigma_h}) > \tau$ represents valid circles.

Questions

Generalized Hough transform


- We want to find a shape defined by its boundary points and a reference point



D. Ballard, [Generalizing the Hough Transform to Detect Arbitrary Shapes](#), Pattern Recognition 13(2), 1981, pp. 111-122.

Generalized Hough transform

- We want to find a shape defined by its boundary points and a reference point
- For every boundary point p , we can compute the displacement vector $r = a - p$ as a function of gradient orientation θ




D. Ballard, [Generalizing the Hough Transform to Detect Arbitrary Shapes](#), Pattern Recognition 13(2), 1981, pp. 111-122.

Generalized Hough transform


- For model shape: construct a table indexed by ϑ storing displacement vectors r as function of gradient direction
- Detection: For each edge point p with gradient orientation ϑ :
 - Retrieve all r indexed with ϑ
 - For each $r(\vartheta)$, put a vote in the Hough space at $p + r(\vartheta)$
- Peak in this Hough space is reference point with most supporting edges
- Assumption: translation is the only transformation here, i.e., orientation and scale are fixed

Example




model shape

Example




displacement vectors for model points

Example




range of voting locations for test point

A diagram of a closed system with five boundaries. The boundaries are: a vertical line on the left, a horizontal line at the top, a horizontal line at the top right, a vertical line on the right, and a diagonal line at the bottom right. Arrows indicate mass flow: an arrow points right from the left vertical boundary; an arrow points down from the top horizontal boundary; an arrow points down from the top right horizontal boundary; an arrow points left from the right vertical boundary; and an arrow points down-left from the bottom right diagonal boundary.




range of voting locations for test point

Example




Example




displacement vectors for model points

Example




range of voting locations for test point

Example




Application in recognition

- Instead of indexing displacements by gradient orientation, index by “visual codeword”



training image



visual codeword with
displacement vectors

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

Application in recognition

- Instead of indexing displacements by gradient orientation, index by “visual codeword”



test image

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

Homography

- The transformation between two views of a planar surface



The transformation between images from two cameras that share the same center



Fitting a homography

- Recall: homogenous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogenous
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogenous
image coordinates

Fitting a homography

- Recall: homogenous coordinates

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Converting *to* homogenous
image coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Converting *from* homogenous
image coordinates

- Equation for homography:

$$\lambda \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Fitting a homography

- Equation for homography:

$$\lambda \begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \quad \lambda \mathbf{x}'_i = \mathbf{H} \mathbf{x}_i = \begin{bmatrix} \mathbf{h}_1^T \\ \mathbf{h}_2^T \\ \mathbf{h}_3^T \end{bmatrix} \mathbf{x}_i$$

9 entries, 8 degrees of freedom
(scale is arbitrary)

$$\mathbf{x}'_i \times \mathbf{H} \mathbf{x}_i = 0$$

$$\mathbf{x}'_i \times \mathbf{H} \mathbf{x}_i = \begin{bmatrix} y'_i \mathbf{h}_3^T \mathbf{x}_i - \mathbf{h}_2^T \mathbf{x}_i \\ \mathbf{h}_1^T \mathbf{x}_i - x'_i \mathbf{h}_3^T \mathbf{x}_i \\ x'_i \mathbf{h}_2^T \mathbf{x}_i - y'_i \mathbf{h}_1^T \mathbf{x}_i \end{bmatrix}$$

$$\begin{bmatrix} 0^T & -\mathbf{x}_i^T & y'_i \mathbf{x}_i^T \\ \mathbf{x}_i^T & 0^T & -x'_i \mathbf{x}_i^T \\ -y'_i \mathbf{x}_i^T & x'_i \mathbf{x}_i^T & 0^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0$$




3 equations, only 2 linearly independent

Direct linear transform

$$\begin{bmatrix} 0^T & \mathbf{x}_1^T & -y'_1 \mathbf{x}_1^T \\ \mathbf{x}_1^T & 0^T & -x'_1 \mathbf{x}_1^T \\ \dots & \dots & \dots \\ 0^T & \mathbf{x}_n^T & -y'_n \mathbf{x}_n^T \\ \mathbf{x}_n^T & 0^T & -x'_n \mathbf{x}_n^T \end{bmatrix} \begin{pmatrix} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \end{pmatrix} = 0 \quad \mathbf{A} \mathbf{h} = 0$$

- H has 8 degrees of freedom (9 parameters, but scale is arbitrary)
- One match gives us two linearly independent equations
- Four matches needed for a minimal solution (null space of 8x9 matrix)
- More than four: homogeneous least squares

Application: Panorama stitching



Recognizing panoramas

- Given contents of a camera memory card, automatically figure out which pictures go together and stitch them together into panoramas



M. Brown and D. Lowe, *["Recognizing Panoramas,"](http://www.cs.ubc.ca/~mbrown/panorama/panorama.html)* ICCV 2003.

<http://www.cs.ubc.ca/~mbrown/panorama/panorama.html>